

# Relational Algebra

# Relational Algebra

- Procedural language
- Six basic operators
  - select:  $\sigma$
  - project:  $\Pi$
  - rename:  $\rho$
  - union:  $\cup$
  - set difference:  $-$
  - Cartesian product:  $\times$
- The operators take one or two relations as inputs and produce a new relation as a result.

# Select Operation – Example

□ Relation r

A	B	C	D
$\alpha$	$\alpha$	1	7
$\alpha$	$\beta$	5	7
$\beta$	$\beta$	12	3
$\beta$	$\beta$	23	10

■  $\sigma_{A=B \wedge D > 5}(r)$

A	B	C	D
$\alpha$	$\alpha$	1	7
$\beta$	$\beta$	23	10

# Select Operation

- Notation:  $\sigma_p(r)$
- $p$  is called the **selection predicate**
- Defined as:

$$\sigma_p(r) = \{t \mid t \in r \text{ and } p(t)\}$$

Where  $p$  is a formula in propositional calculus consisting of **terms** connected by :  $\wedge$  (**and**),  $\vee$  (**or**),  $\neg$  (**not**)  
Each **term** is one of:

$\langle \text{attribute} \rangle \text{ } op \text{ } \langle \text{attribute} \rangle$  or  $\langle \text{constant} \rangle$

where  $op$  is one of:  $=, \neq, >, \geq, <, \leq$

- Example of selection:

$$\sigma_{dept\_name="Physics"}(instructor)$$

# Project Operation – Example

○ Relation  $r$ :

$A$	$B$	$C$
$\alpha$	10	1
$\alpha$	20	1
$\beta$	30	1
$\beta$	40	2

□  $\Pi_{A,C}(r)$

$A$	$C$
$\alpha$	1
$\alpha$	1
$\beta$	1
$\beta$	2

 $=$ 

$A$	$C$
$\alpha$	1
$\beta$	1
$\beta$	2

# Project Operation

○ Notation:  $\Pi_{A_1, A_2, \dots, A_k}(r)$

where  $A_1, A_2$  are attribute names and  $r$  is a relation name.

- The result is defined as the relation of  $k$  columns obtained by erasing the columns that are not listed
- Duplicate rows removed from result, since relations are sets
- Example: To eliminate the *dept\_name* attribute of *instructor*

$$\Pi_{ID, name, salary}(instructor)$$

# Union Operation – Example

○ Relations  $r, s$ :

$A$	$B$
$\alpha$	1
$\alpha$	2
$\beta$	1

$r$

$A$	$B$
$\alpha$	2
$\beta$	3

$s$

□  $r \cup s$ :

$A$	$B$
$\alpha$	1
$\alpha$	2
$\beta$	1
$\beta$	3

# Union Operation

- Notation:  $r \cup s$
- Defined as:

$$r \cup s = \{t \mid t \in r \text{ or } t \in s\}$$

- For  $r \cup s$  to be valid.
  1.  $r, s$  must have the **same arity** (same number of attributes)
  2. The attribute domains must be **compatible** (example: 2<sup>nd</sup> column of  $r$  deals with the same type of values as does the 2<sup>nd</sup> column of  $s$ )
- Example: to find all courses taught in the Fall 2009 semester, or in the Spring 2010 semester, or in both

$$\Pi_{course\_id} (\sigma_{semester="Fall" \wedge year=2009} (section)) \cup \\ \Pi_{course\_id} (\sigma_{semester="Spring" \wedge year=2010} (section))$$



# Set difference of two relations

- Relations  $r, s$ :

$A$	$B$
$\alpha$	1
$\alpha$	2
$\beta$	1

$r$

$A$	$B$
$\alpha$	2
$\beta$	3

$s$

□  $r - s$ :

$A$	$B$
$\alpha$	1
$\beta$	1

# Set Difference Operation

- Notation  $r - s$
- Defined as:

$$r - s = \{t \mid t \in r \text{ and } t \notin s\}$$

- Set differences must be taken between **compatible** relations.
  - $r$  and  $s$  must have the **same** arity
  - attribute domains of  $r$  and  $s$  must be compatible
- Example: to find all courses taught in the Fall 2009 semester, but not in the Spring 2010 semester

$$\Pi_{course\_id} (\sigma_{semester="Fall" \wedge year=2009} (section)) - \Pi_{course\_id} (\sigma_{semester="Spring" \wedge year=2010} (section))$$

# Cartesian-Product Operation–Example

□ Relations  $r, s$ :

$A$	$B$
$\alpha$	1
$\beta$	2

$r$

$C$	$D$	$E$
$\alpha$	10	a
$\beta$	10	a
$\beta$	20	b
$\gamma$	10	b

$s$

□  $r \times s$ :

$A$	$B$	$C$	$D$	$E$
$\alpha$	1	$\alpha$	10	a
$\alpha$	1	$\beta$	10	a
$\alpha$	1	$\beta$	20	b
$\alpha$	1	$\gamma$	10	b
$\beta$	2	$\alpha$	10	a
$\beta$	2	$\beta$	10	a
$\beta$	2	$\beta$	20	b
$\beta$	2	$\gamma$	10	b

# Cartesian-Product Operation

- Notation  $r \times s$

- Defined as:

$$r \times s = \{t \ q \mid t \in r \textbf{ and } q \in s\}$$

- Assume that attributes of  $r(R)$  and  $s(S)$  are disjoint. (That is,  $R \cap S = \emptyset$ ).

- If attributes of  $r(R)$  and  $s(S)$  are not disjoint, then renaming must be used because the same name is used in both relations.

# Composition of Operations

- Can build expressions using multiple operations
- Example:  $\sigma_{A=C}(r \times s)$
- $r \times s$

A	B	C	D	E
$\alpha$	1	$\alpha$	10	a
$\alpha$	1	$\beta$	10	a
$\alpha$	1	$\beta$	20	b
$\alpha$	1	$\gamma$	10	b
$\beta$	2	$\alpha$	10	a
$\beta$	2	$\beta$	10	a
$\beta$	2	$\beta$	20	b
$\beta$	2	$\gamma$	10	b

- $\sigma_{A=C}(r \times s)$

A	B	C	D	E
$\alpha$	1	$\alpha$	10	a
$\beta$	2	$\beta$	10	a
$\beta$	2	$\beta$	20	b

# Rename Operation

- Allows us to name, and therefore to refer to, the results of relational-algebra expressions.
- Allows us to refer to a relation by more than one name.
- Example:

$$\rho_x(E)$$

returns the expression  $E$  under the name  $X$

- If a relational-algebra expression  $E$  has arity  $n$ , then

$$\rho_{x(A_1, A_2, \dots, A_n)}(E)$$

returns the result of expression  $E$  under the name  $X$ , and with the attributes renamed to  $A_1, A_2, \dots, A_n$ .

# Example Query

- Find the largest salary in the university
  - Step 1: find instructor salaries that are less than some other instructor salary (i.e. not maximum)
    - using a copy of *instructor* under a new name *d*
    - $\Pi_{instructor.salary} (\sigma_{instructor.salary < d.salary} (instructor \times \rho_d (instructor)))$
  - Step 2: Find the largest salary
    - $\Pi_{salary} (instructor) - \Pi_{instructor.salary} (\sigma_{instructor.salary < d.salary} (instructor \times \rho_d (instructor)))$

Salary
500
800
900
1000
700

Salary
500
800
900
700

Salary
1000

# Example Queries

- Find the names of all instructors in the Physics department, along with the *course\_id* of all courses they have taught

- Query 1

$$\Pi_{instructor.ID, course\_id} (\sigma_{dept\_name="Physics"} (\sigma_{instructor.ID=teaches.ID} (instructor \times teaches)))$$

- Query 2

$$\Pi_{instructor.ID, course\_id} (\sigma_{instructor.ID=teaches.ID} (\sigma_{dept\_name="Physics"} (instructor \times teaches)))$$



# Formal Definition

○ Let  $E_1$  and  $E_2$  be relational-algebra expressions; the following are all relational-algebra expressions:

- $E_1 \cup E_2$
- $E_1 - E_2$
- $E_1 \times E_2$
- $\sigma_p(E_1)$ ,  $P$  is a predicate on attributes in  $E_1$
- $\Pi_s(E_1)$ ,  $S$  is a list consisting of some of the attributes in  $E_1$
- $\rho_x(E_1)$ ,  $x$  is the new name for the result of  $E_1$


# Additional Operations

We define additional operations that do not add any power to the relational algebra, but that simplify common queries.

- Set intersection
- Natural join
- Assignment
- Outer join

# Set-Intersection Operation

- Notation:  $r \cap s$
- Defined as:
  - $r \cap s = \{ t \mid t \in r \textbf{ and } t \in s \}$
- Assume:
  - $r, s$  have the *same arity*
  - attributes of  $r$  and  $s$  are compatible
- Note:  $r \cap s = r - (r - s)$



# Set-Intersection Operation – Example

○ Relation  $r, s$ :

$A$	$B$
$\alpha$	1
$\alpha$	2
$\beta$	1

$r$

$A$	$B$
$\alpha$	2
$\beta$	3

$s$

○  $r \cap s$

$A$	$B$
$\alpha$	2

# Natural-Join Operation

- Notation:  $r \bowtie s$
- Let  $r$  and  $s$  be relations on schemas  $R$  and  $S$  respectively.  
Then,  $r \bowtie s$  is a relation on schema  $R \cup S$  obtained as follows:
  - Consider each pair of tuples  $t_r$  from  $r$  and  $t_s$  from  $s$ .
  - If  $t_r$  and  $t_s$  have the same value on each of the attributes in  $R \cap S$ , add a tuple  $t$  to the result, where
    - $t$  has the same value as  $t_r$  on  $r$
    - $t$  has the same value as  $t_s$  on  $s$
- Example:
  - $R = (A, B, C, D)$
  - $S = (E, B, D)$
  - Result schema =  $(A, B, C, D, E)$
  - $r \bowtie s$  is defined as:  
$$\Pi_{r.A, r.B, r.C, r.D, s.E} (\sigma_{r.B = s.B \wedge r.D = s.D} (r \times s))$$

# Natural Join Example

○ Relations  $r, s$ :

A	B	C	D
$\alpha$	1	$\alpha$	a
$\beta$	2	$\gamma$	a
$\gamma$	4	$\beta$	b
$\alpha$	1	$\gamma$	a
$\delta$	2	$\beta$	b

$r$

B	D	E
1	a	$\alpha$
3	a	$\beta$
1	a	$\gamma$
2	b	$\delta$
3	b	$\epsilon$

$s$

□  $r \bowtie s$

A	B	C	D	E
$\alpha$	1	$\alpha$	a	$\alpha$
$\alpha$	1	$\alpha$	a	$\gamma$
$\alpha$	1	$\gamma$	a	$\alpha$
$\alpha$	1	$\gamma$	a	$\gamma$
$\delta$	2	$\beta$	b	$\delta$

# Natural Join

- Find the names of all instructors in the Comp. Sci. department together with the course titles of all the courses that the instructors teach

- $\Pi_{name, title} (\sigma_{dept\_name="Comp. Sci."} (instructor \bowtie teaches \bowtie course))$

- Natural join is associative

- $(instructor \bowtie teaches) \bowtie course$  is equivalent to  $instructor \bowtie (teaches \bowtie course)$

- Natural join is commutative

- $instruct \bowtie teaches$  is equivalent to  $teaches \bowtie instructor$

# Assignment Operation

- The assignment operation ( $\leftarrow$ ) provides a convenient way to express complex queries.
  - Write query as a sequential program consisting of
    - a series of assignments
    - followed by an expression whose value is displayed as a result of the query.
  - Assignment must always be made to a temporary relation variable.
  - Example:

$$\begin{aligned}temp1 &\leftarrow \Pi_{R-S}(r) \\temp2 &\leftarrow \Pi_{R-S}((temp1 \times s) - \Pi_{R-S,S}(r)) \\result &= temp1 - temp2\end{aligned}$$



# Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values:
  - *null* signifies that the value is unknown or does not exist
  - All comparisons involving *null* are **false** by definition.

# Outer Join – Example

- Relation *instructor1*

<i>ID</i>	<i>name</i>	<i>dept_name</i>
10101	Srinivasan	Comp. Sci.
12121	Wu	Finance
15151	Mozart	Music

- Relation *teaches1*

<i>ID</i>	<i>course_id</i>
10101	CS-101
12121	FIN-201
76766	BIO-101

# Outer Join – Example

- Join

*instructor* ⋈ *teaches*

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>course_id</i>
10101	Srinivasan	Comp. Sci.	CS-101
12121	Wu	Finance	FIN-201

- Left Outer Join

*instructor* ⋈<sub>L</sub> *teaches*

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>course_id</i>
10101	Srinivasan	Comp. Sci.	CS-101
12121	Wu	Finance	FIN-201
15151	Mozart	Music	<i>null</i>

# Outer Join – Example

- Right Outer Join

*instructor* ⋈<sub>⊃</sub> *teaches*

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>course_id</i>
10101	Srinivasan	Comp. Sci.	CS-101
12121	Wu	Finance	FIN-201
76766	null	null	BIO-101

- Full Outer Join

*instructor* ⋈<sub>⊃</sub> *teaches*

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>course_id</i>
10101	Srinivasan	Comp. Sci.	CS-101
12121	Wu	Finance	FIN-201
15151	Mozart	Music	null
76766	null	null	BIO-101

# Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*.
- Aggregate functions simply ignore null values (as in SQL)
- For duplicate elimination and grouping, null is treated like any other value, and two nulls are assumed to be the same (as in SQL)

# Null Values

- Comparisons with null values return the special truth value: *unknown*
- Three-valued logic using the truth value *unknown*:
  - OR: (*unknown or true*) = *true*,  
(*unknown or false*) = *unknown*  
(*unknown or unknown*) = *unknown*
  - AND: (*true and unknown*) = *unknown*,  
(*false and unknown*) = *false*,  
(*unknown and unknown*) = *unknown*
  - NOT: (**not** *unknown*) = *unknown*
  - In SQL “*P is unknown*” evaluates to true if predicate *P* evaluates to *unknown*
- Result of select predicate is treated as *false* if it evaluates to *unknown*

# Division Operator

- Given relations  $r(R)$  and  $s(S)$ , such that  $S \subset R$ ,  $r \div s$  is the largest relation  $t(R-S)$  such that  $t \times s \subseteq r$
- E.g. let  $r(ID, course\_id) = \Pi_{ID, course\_id}(takes)$  and  $s(course\_id) = \Pi_{course\_id}(\sigma_{dept\_name="Biology"}(course))$

then  $r \div s$  gives us **students who have taken all courses in the Biology department**

- Can write  $r \div s$  as

$$temp1 \leftarrow \Pi_{R-S}(r)$$

$$temp2 \leftarrow \Pi_{R-S}((temp1 \times s) - \Pi_{R-S,S}(r))$$

$$result = temp1 - temp2$$

- The result to the right of the  $\leftarrow$  is assigned to the relation variable on the left of the  $\leftarrow$ .
- May use variable in subsequent expressions.

# Not Exists (from the SQL slides)

- Find **all** students who have taken **all** courses offered in the Biology department.

```
select distinct S.ID, S.name  
from student as S  
where not exists ( (select course_id  
                    from course  
                    where dept_name = 'Biology')  
except  
                  (select T.course_id  
                   from takes as T  
                   where S.ID = T.ID));
```

3- takes each student and tests whether the set of all courses that the student has taken contains the set of all courses offered in the Biology department.

1- finds the set of all courses offered in the Biology department

2- finds all the courses that student S.ID has taken

□ Note that  $X - Y = \emptyset \Leftrightarrow X \subseteq Y$



# Extended Relational-Algebra-Operations

- Generalized Projection
- Aggregate Functions

# Generalized Projection

- Extends the projection operation by allowing arithmetic functions to be used in the projection list.

$$\Pi_{F_1, F_2, \dots, F_n}(E)$$

$E$  is any relational-algebra expression

- Each of  $F_1, F_2, \dots, F_n$  are arithmetic expressions involving constants and attributes in the schema of  $E$ .
- Given relation *instructor*(*ID*, *name*, *dept\_name*, *salary*) where *salary* is annual salary, get the same information but with monthly salary

$$\Pi_{ID, name, dept\_name, salary/12}(instructor)$$

# Aggregate Functions and Operations

- **Aggregation function** takes a collection of values and returns a single value as a result.

**avg**: average value

**min**: minimum value

**max**: maximum value

**sum**: sum of values

**count**: number of values

- **Aggregate operation** in relational algebra

$$G_1, G_2, \dots, G_n \quad \mathcal{G} \quad F_1(A_1), F_2(A_2), \dots, F_n(A_n) (E)$$

$E$  is any relational-algebra expression

- $G_1, G_2, \dots, G_n$  is a list of attributes on which to group (can be empty)
- Each  $F_i$  is an aggregate function
- Each  $A_i$  is an attribute name

# Aggregate Operation – Example

○ Relation  $r$ :

$A$	$B$	$C$
$\alpha$	$\alpha$	7
$\alpha$	$\beta$	7
$\beta$	$\beta$	3
$\beta$	$\beta$	10

□  $G_{\text{sum}(c)}(r)$

<b>sum(c)</b>
27

# Aggregate Operation – Example

- Find the average salary in each department

*dept\_name*  $\mathcal{G}$  *avg(salary)* (*instructor*)

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

<i>dept_name</i>	<i>avg_salary</i>
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000

# Aggregate Functions (Cont.)

- Result of aggregation does not have a name
  - Can use rename operation to give it a name
  - For convenience, we permit renaming as part of aggregate operation

*dept\_name* **G** **avg**(salary) **as** *avg\_sal* (*instructor*)

# SQL and Relational Algebra

- **select**  $A1, A2, \dots, An$   
**from**  $r1, r2, \dots, rm$   
**where**  $P$

is equivalent to the following expression in multiset relational algebra

$$\Pi_{A1, \dots, An} (\sigma_P (r1 \times r2 \times \dots \times rm))$$

- **select**  $A1, A2, \text{sum}(A3)$   
**from**  $r1, r2, \dots, rm$   
**where**  $P$   
**group by**  $A1, A2$

is equivalent to the following expression in multiset relational algebra

$$\Pi_{A1, A2} G_{\text{sum}(A3)} (\sigma_P (r1 \times r2 \times \dots \times rm))$$

# Modification of the Database

- The content of the database may be modified using the following operations:
  - Deletion
  - Insertion
  - Updating
- All these operations can be expressed using the assignment operator



# Deletion

- A delete request is expressed similarly to a query, except instead of displaying tuples to the user, the selected tuples are removed from the database.
- Can delete only whole tuples; cannot delete values on only particular attributes
- A deletion is expressed in relational algebra by:

$$r \leftarrow r - E$$

where  $r$  is a relation and  $E$  is a relational algebra query.

# Deletion Examples

- Delete all account records in the Perry branch.

$account \leftarrow account - \sigma_{branch\_name = "Perry"}(account)$

- Delete all loan records with amount in the range of 0 to 50

$loan \leftarrow loan - \sigma_{amount \geq 0 \text{ and } amount \leq 50}(loan)$

- Delete all accounts at branches located in Needham.

$r_1 \leftarrow \sigma_{branch\_city = "Needham"}(account \bowtie branch)$

$r_2 \leftarrow \Pi_{account\_number, branch\_name, balance}(r_1)$

$r_3 \leftarrow \Pi_{customer\_name, account\_number}(r_2 \bowtie depositor)$

$account \leftarrow account - r_2$

$depositor \leftarrow depositor - r_3$

Information of the accounts at the branches located in Needham

Filter the values of the attributes required for the schema *account*

Combine that data with *depositor* and select the attributes of the schema *depositor*

Remove selected data from account and from depositor

# Insertion

- To insert data into a relation, we either:
  - specify a tuple to be inserted
  - write a query whose result is a set of tuples to be inserted
- in relational algebra, an insertion is expressed by:

$$r \leftarrow r \cup E$$

where  $r$  is a relation and  $E$  is a relational algebra expression.

- The insertion of a single tuple is expressed by letting  $E$  be a constant relation containing one tuple.

# Insertion Examples

- Insert information in the database specifying that Smith has \$1200 in account A-973 at the Perry branch.

$account \leftarrow account \cup \{("A-973", "Perry", 1200)\}$

$depositor \leftarrow depositor \cup \{("Smith", "A-973")\}$

- Provide as a gift for all loan customers in the Perry branch, a \$200 savings account. Let the loan number serve as the account number for the new savings account.

$r_1 \leftarrow (\sigma_{branch\_name = "Perryridge"}(borrower \bowtie loan))$

$account \leftarrow account \cup \Pi_{loan\_number, branch\_name, 200}(r_1)$

$depositor \leftarrow depositor \cup \Pi_{customer\_name, loan\_number}(r_1)$

Information of the loans and customers at the branches located in Perry

Filter the values of the attributes required for the schema *account* and add them to the relation *account*

Filter the values of the attributes required for the schema *depositor* and add them to the relation *depositor*

# Updating

- A mechanism to change a value in a tuple
- Use the generalized projection operator to do this task

$$r \leftarrow \Pi_{F_1, F_2, \dots, F_I}(r)$$

- Each  $F_i$  is either
  - the  $i^{\text{th}}$  attribute of  $r$ , if the  $i^{\text{th}}$  attribute is not updated, or,
  - if the attribute is to be updated  $F_i$  is an expression, involving only constants and the attributes of  $r$ , which gives the new value for the attribute

# Update Examples

- Make interest payments by increasing all balances by 5 percent.

$account \leftarrow \Pi_{account\_number, branch\_name, balance * 1.05} (account)$

- Pay all accounts with balances over \$10,000 6 percent interest and pay all others 5 percent

$account \leftarrow \Pi_{account\_number, branch\_name, balance * 1.06} (\sigma_{BAL > 10000} (account)) \cup \Pi_{account\_number, branch\_name, balance * 1.05} (\sigma_{BAL \leq 10000} (account))$

# Example Queries

- Find the names of all customers who have a loan and an account at bank.

# Example Queries

- Find the name of all customers who have a loan at the bank and the loan amount



## Example Queries

- Find all customers who have an account from at least the “Downtown” and the Uptown” branches.

# Bank Example Queries

- Find all customers who have an account at all branches located in Brooklyn city.